

Shift-Reduce CCG Parsing

Yue Zhang and Stephen Clark
University of Cambridge

Outline

- Introduction
- CCG
- Our Shift-Reduce Parser
- Experiments

Outline

- Introduction
- CCG
- Our Shift-Reduce Parser
- Experiments

Introduction

- Combinatory Categorical Grammar (CCG)
 - Lexicalized grammar
 - Successfully applied to a range of problems:
 - treebank creation (Hockenmaier and Steedman, 2007)
 - syntactic parsing (Hockenmaier, 2003; Clark and Curran, 2007)
 - logical form construction (Bos et al., 2004)
 - surface realization (White and Rajkumar, 2009)

Introduction

- Combinatory Categorical Grammar (CCG)
 - Binary branching
 - Naturally compatible with bottom-up parsing algorithms, such as shift-reduce and CKY (Ades and Steedman, 1982; Steedman, 2000)
 - Previous work only considered chart-parsing (Clark and Curran, 2007; Hockenmaier, 2003; Fowler and Penn, 2010)
 - We fill a gap in the literature by developing a shift-reduce parser for CCG

Introduction

- Shift-reduce parsing
 - Popular for dependency parsing, following the work of Yamada and Matsumoto (2003) and Nivre and Scholz (2004)
 - Uses a stack and a set of shift-reduce actions to build parses
 - Compared to chart parsing:
 - allows highly efficient parsing
 - allows definition of rich features
 - makes different errors
 - Leads to higher accuracy by combination with chart-parsing

Introduction

- Our shift-reduce parser vs. chart-based C&C
 - C&C (Clark and Curran, 2007) is a chart-based CCG parser
 - Both use global discriminative models
 - Heuristic beam-search with rich features vs. optimal dynamic-programming with restricted features
 - 85.53% F-score vs. 85.45% overall
 - Higher precision and lower recall
 - Naturally handles cases when spanning analysis cannot be found

Outline

- Introduction
- CCG
- Our Shift-Reduce Parser
- Experiments

CCG Parsing

- Lexical categories
 - basic categories: N (nouns), NP (noun phrases), PP (prepositional phrases), ...
 - complex categories: $S \backslash NP$ (intransitive verbs), $(S \backslash NP) / NP$ (transitive verbs), ...
- Adjacent phrases are combined to form larger phrases using category combination e.g.:
 - function application: $NP \ S \backslash NP \Rightarrow S$
 - function composition: $(S \backslash NP) / (S \backslash NP) \ (S \backslash NP) / NP \Rightarrow (S \backslash NP) / NP$
- Unary rules change the type of a phrase
 - Type raising: $NP \Rightarrow S / (S \backslash NP)$
 - Type changing: $S[pss] \backslash NP \Rightarrow NP \backslash NP$

CCG Parsing

- An example derivation

IBM bought

Lotus

CCG Parsing

- An example derivation

IBM	bought	Lotus
NP	$(S[dcl] \backslash NP) / NP$	NP

CCG Parsing

- An example derivation

IBM	bought	Lotus
NP	$(S[dcl] \backslash NP) / NP$	NP
<hr/>		
	$S[dcl] \backslash NP$	

CCG Parsing

- An example derivation

IBM	bought	Lotus
NP	$(S[dcl] \backslash NP) / NP$	NP

$S[dcl] \backslash NP$

$S[dcl]$

CCG Parsing


- Rule extraction
 - Manually define the combinatory rule schemas (Steedman, 2000; Clark and Curran, 2007)
 - Schema: $X/Y \ Y \Rightarrow X$
 - Instance: $(S \backslash NP) / NP \ NP \Rightarrow S \backslash NP$
 - Extracting rule instances from corpus (Hockenmaier, 2003; Fowler and Penn, 2010)

Outline

- Introduction
- CCG
- Our Shift-Reduce Parser
- Experiments



The Shift-Reduce Parser

- The shift-reduce parser
 - A *stack* of partial derivations
 - A *queue* of input words
 - A set of shift-reduce *actions* to build parsers
 - *SHIFT*
 - *COMBINE*
 - *UNARY*
 - *FINISH*





The Shift-Reduce Parser

- Shift-reduce actions
 - *SHIFT-X*
 - Pushes the head of the queue onto the stack
 - Assigns label X (a lexical category)
 - *SHIFT* action performs lexical category disambiguation





The Shift-Reduce Parser

- Shift-reduce actions
 - *COMBINE-X*
 - Pops the top two nodes off the stack
 - Combines into a new node X , and push it onto stack
 - Corresponds to the use of a combinatory rule in CCG



The Shift-Reduce Parser

- Shift-reduce actions
 - *UNARY-X*
 - Pops the top of the stack
 - Create a new node with category *X*; pushes it onto stack
 - Corresponds to the use of a unary rule in CCG



The Shift-Reduce Parser

- Shift-reduce actions
 - *FINISH*
 - Terminates the parsing process
 - Can be applied when all input words have been pushed onto the stack
 - Allows fragmentary analysis:
 - when the stack holds multiple items that cannot be combined
 - such cases can arise from incorrect lexical category assignment

The Shift-Reduce Parser

- An example parsing process



IBM bought Lotus yesterday

initial

The Shift-Reduce Parser

- An example parsing process



bought Lotus yesterday

SHIFT

The Shift-Reduce Parser

- An example parsing process

$NP_{IBM} ((S[dc] \backslash NP) / NP)_{bought}$

Lotus yesterday

SHIFT

The Shift-Reduce Parser


- An example parsing process

$NP_{IBM} ((S[dcI]NP)/NP)_{bought} NP_{Lotus}$	yesterday
--	-----------

SHIFT

The Shift-Reduce Parser


- An example parsing process



COMBINE

The Shift-Reduce Parser


- An example parsing process



SHIFT

The Shift-Reduce Parser


- An example parsing process



COMBINE

The Shift-Reduce Parser

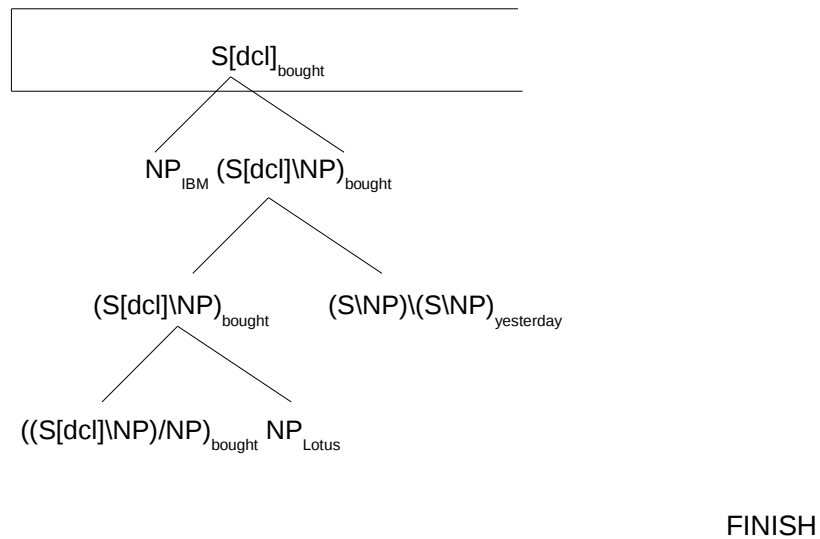
- An example parsing process



COMBINE

The Shift-Reduce Parser

- An example parsing process



Decoding

- Candidate item $\langle S, Q, F \rangle$
 - S – stack
 - Q – queue
 - F – boolean: parsing finished (finish action applied)
 - Start item: $\langle \text{empty}, \text{input}, \text{false} \rangle$
- A derivation is built:
 - from the start item
 - by repeated application of actions
 - until the item is finished

Decoding


- Beam-search decoding
 - An *agenda* holds N -best unfinished items at each step ($N=16$)
 - A *candidate output* records the best finished item currently found
 - Initially the agenda contains only the start item
 - At each step
 - Each item in the agenda is expanded with all actions
 - Agenda replaced with N -best new unfinished items
 - Candidate output is updated with new finished items
 - Repeat until the agenda is empty
 - Candidate output is taken as the final output

Model and Training

- Global linear model
- Trained with:
 - the averaged perceptron (Collins, 2002)
 - “early-update” strategy (Collins and Roark, 2004)
- Zhang and Clark (2008), Zhang and Clark (2009), Huang et al., (2009)

Features

context



- Stack nodes: $S_0 S_1 S_2 S_3$
- Queue nodes: $Q_0 Q_1 Q_2 Q_3$
- Stack subnodes: $S_{0L} S_{0R} S_{0U} S_{1L/R/U}$

S0wp, S0c, S0pc, S0wc,
S1wp, S1c, S1pc, S1wc,
S2pc, S2wc,
S3pc, S3wc,

Q0wp, Q1wp, Q2wp, Q3wp,

S0Lpc, S0Lwc, S0Rpc, S0Rwc,
S0Upc, S0Uwc,
S1Lpc, S1Lwc, S1Rpc, S1Rwc,
S1Upc, S1Uwc,

S0wcS1wc, S0cS1w, S0wS1c, S0cS1c,
S0wcQ0wp, S0cQ0wp, S0wcQ0p, S0cQ0p,
S1wcQ0wp, S1cQ0wp, S1wcQ0p, S1cQ0p,

S0wcS1cQ0p, S0cS1wcQ0p, S0cS1cQ0wp,
S0cS1cQ0p, S0pS1pQ0p,
S0wcQ0pQ1p, S0cQ0wpQ1p, S0cQ0pQ1wp,
S0cQ0pQ1p, S0pQ0pQ1p,
S0wcS1cS2c, S0cS1wcS2c, S0cS1cS2wc,
S0cS1cS2c, S0pS1pS2p,

S0cS0HcS0Lc, S0cS0HcS0Rc,
S1cS1HcS1Rc,
S0cS0RcQ0p, S0cS0RcQ0w,
S0cS0LcS1c, S0cS0LcS1w,
S0cS1cS1Rc, S0wS1cS1Rc.

Outline

- Introduction
- CCG
- Our Shift-Reduce Parser
- Experiments

Experimental Data

- CCGBank (Hockenmaier and Steedman, 2007)
- Split into three subsets:
 - Training (section 02 – 21)
 - Development (section 00)
 - Testing (section 23)
- Extract CCG rules
 - Binary instances: 3070
 - Unary instances: 191

Evaluation

- Standard evaluation: F-scores over labeled CCG dependencies:
 - use evaluate script from C&C tools
- Our parser produces CCG derivation trees:
 - use generate script from C&C tools to transform derivations into dependencies

POS-tagging and Supertagging

- Supertagging is the process of lexical category assignment
- Use pos and msuper from C&C tools
- 10-fold cross-validation for training data POS-tagging and supertagging
- Multiple supertags (lexical categories) per word
 - Use a probability threshold: β
 - A small β value of 0.0001 leads to comparatively more supertags per word
 - Average 5.4 per word on the development data

Development Test Accuracies

- Labeled precision (lp.), recall (lr.) and F-score (lf.)
- Complete sentence matches (lsent.)
- Lexical category accuracy (cats.)
- Measured on:
 - all input sentences
 - C&C coverage


	lp.	lr.	lf.	lsent.	cats.	evaluated on
shift-reduce	87.15%	82.95%	85.00%	33.82%	92.77%	all sentences
C&C (normal-form)	85.22%	82.52%	83.85%	31.63%	92.40%	all sentences
shift-reduce	87.55%	83.63%	85.54%	34.14%	93.11%	99.06% (C&C coverage)
C&C (hybrid)	--	--	85.25%	--	--	99.06% (C&C coverage)
C&C (normal-form)	85.22%	84.29%	84.76%	31.93%	92.83%	99.06% (C&C coverage)

Error Comparisons

- Two different solutions to the same problem
 - Approximate beam-search with rich features
 - Optimal dynamic-programming with restricted features
- Similar to comparison between shift-reduce based MaltParser (Nivre et al., 2006) and chart-based MSTParser (McDonald et al., 2006)
- Characterize errors by sentence size, dependency length and type

Error Comparisons

- As sentence length increases
 - Both parsers give lower performance
 - No difference in the rate of accuracy degradation
- When dependency length increases



Error Comparisons

- On dependency types

category	arg	lp. (ours)	lp. (C&C)	lr. (ours)	lr. (C&C)	lf. (ours)	lf. (C&C)	freq.
N/N	1	95.77	95.28	95.79	95.62	95.78	95.45	7288
NP/N	1	96.70	96.57	96.59	96.03	96.65	96.30	4101
(NP\NP)/NP	2	83.19	82.17	89.24	88.90	86.11	85.40	2379
(NP\NP)/NP	1	82.53	81.58	87.99	85.74	85.17	83.61	2174
((S\NP)\ (S\NP))/NP	3	77.60	71.94	71.58	73.32	74.47	72.63	1147
((S\NP)\ (S\NP))/NP	2	76.30	70.92	70.60	71.93	73.34	71.42	1058
((S[dcI]\NP)/NP	2	85.60	81.57	84.30	86.37	84.95	83.90	917
PP/NP	1	73.76	75.06	72.83	70.09	73.29	72.49	876
((S[dcI]\NP)/NP	1	85.32	81.62	82.00	85.55	83.63	83.54	872
((S\NP)\(S\NP))	2	84.44	86.85	86.60	86.73	85.51	86.79	746

Test results

- F&P = Fowler and Penn (2010)

	lp.	lr.	lf.	lsent.	cats.	evaluated
shift-reduce	87.43	83.61	85.48	35.19	93.12	all sentences
C&C (normal-form)	85.58	82.85	84.20	32.90	92.84	all sentences
shift-reduce	87.43	83.71	85.53	35.34	93.15	99.58% (C&C coverage)
C&C (hybrid)	86.17	84.74	85.45	32.92	92.98	99.58% (C&C coverage)
C&C (normal-form)	85.48	84.60	85.04	33.08	92.86	99.58% (C&C coverage)
F&P (Petrov I-5)*	86.29	85.73	86.01	--	--	-- (F&P \cap C&C coverage; 96.65% on dev. test)
C&C hybrid*	86.46	85.11	85.78	--	--	-- (F&P \cap C&C coverage; 96.65% on dev. test)

Conclusions

- First work to show competitive results for shift-reduce CCG parsing
- Comparison between shift-reduce parsing and chart parsing for CCG
- Complementary errors made by our parser and C&C suggests the possibility of parser combination

Thank you!

(thanks to Michael Auli and Joakim Nivre for
assistance)